

configuration	median latency	90th percentile	mean cost per run	model calls per run
C1 deciding / cloud	4.54 s	8.62 s	\$0.0169	2.37
C2 deciding / on-device	9.08 s	22.69 s	\$0	2.65
C3 routing / cloud	2.55 s	5.91 s	\$0.0092	1.33
C4 routing / on-device	4.88 s	10.58 s	\$0	1.33